

# Backup and Restore databases

---

Saving a copy of data and recovering from data loss

## Scenarios where backup and recovery might be necessary:

- After unplanned shutdown
- Accidental deletion
- Data corruption
- Moving data to a different database system
- Sharing data with business partners
- Using a copy of the data for development or test

## Physical vs Logical backup

- Logical backup contains DDL, DML to recreate
  - Can reclaim wasted space
  - Slow and may impact performance
  - Granular : You can backup individual data or table..
  - Backup/restore, import/export, dump & load utilities
- Physical backup copies physical files, including logs, and configuration
  - smaller and quicker
  - less granular
  - Can only restore to similar RDBMS
  - common for specialized storage and cloud

## What to backup:

- Database
- Schema
- Tables
- Subset of data
- Other objects

Can customize backup policy(what, how frequently,...)

Check if backup is valid, check if restore policy works, ensure that backup files are secure

Other configuration like:

- compression level (more time, less space for higher compression level)
- Encryption (Reduces risk of data being compromised, more time)

## Types of backup:

- Full backup

- backs up all the specified data
- multiple copies means storing many instances of a large file
- Only storing one copy risks data loss if file is corrupt
- Could be needlessly backing up unchanged data
- Most secure (stored offline)
- Point in time backup
  - Uses logged transactions to restore to an earlier point in time.
  - If issue at 12:25 and last full backup was at 11:20, restore the full backup, apply all transactions before that specific transaction.
  - MySQL calls it binary log, PostgreSQL calls it write ahead log, DB2 calls it transaction log.
- Differential backup
  - A copy of data that has changed since the last full backup was taken.
  - Small size.
- Incremental backups
  - A copy of data that has changed since the last backup of any type was taken.

## Backup Policies:

- Hot or online backups:
  - Taken when data is in use
  - User can continue to use db but lower performance
  - Can impact integrity
  - Generally in available server, faster
- Cold or offline backups:
  - Stop the db and then take backup
  - Eliminates data integrity problem
  - Can't be used in 24/7 environment as it impact availability
  - Stored in external server, more time

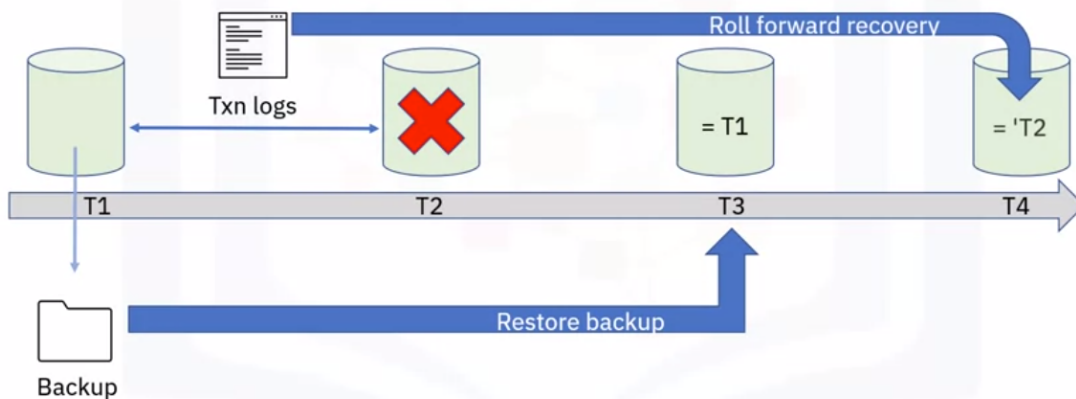
## What should be considered?:

- Frequency:
  - Regularly changing?
  - Size of table
- Schedule:
  - Is data accessed equally across 24-hour day?
  - Is it accessed at weekends?
- Automation
- Backup may need to be enabled in cloud too. (Automatic or Manual)
  - Options depends on provider.
  - Third party solution if dbms doesn't provide it.

## Using DB transaction logs:

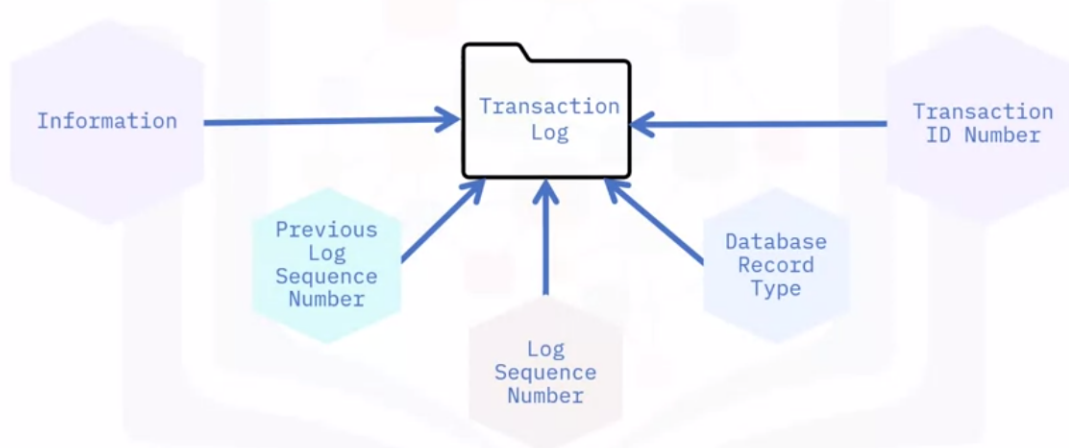
Keep track of transaction that change the db. Can be used to recover.

### Transaction log usage example



- You can change where to store
- Isolate on different volumes from data (recommended)
  - Performance(writes don't conflict) and recoverability
- Log mirroring: Store second copy of log
- Log shipping: automatically send logs somewhere else
- May or maynot be enabled by default

### Basic anatomy of a database log



## Security and User management

### Identifying and mitigating risks

- On premise servers:
  - Who has access?

- What are security measures in place?
- Cloud servers:
  - Providers are responsible (check documentation)
- OS should be secure
  - Regularly patched
  - system hardening (can withstand latest known vulnerability)
  - access monitoring
- RDBMS
  - latest update
  - security features
  - only some have access

User need to be authenticated and allowed access

## Authentication and authorization

- Authentication is how user verifies that they are who they claim to be
  - Username password
  - biometrics lock
  - External tools like login IDs, two factor authentication etc.
- Authorization is granting privileges to users, groups , roles.
  - On table: select, insert, update, delete, alter table.
  - Principle of least privilege: only allow least privelege required.
- Monitoring:
  - who access what
  - what actions they perform
- Auditing:
  - actual access against a security plan
- Encryption:
  - Intruder need to decrypt
- There are regulations:
  - Algorithm
  - Key management
- Consider performance impact too

Even the most secure hardware, RDBMS, and OS can be prone to SQL injections and improper code. So, check those.

## Users, groups and roles

- User is an account.
  - Explicitly need to be granted privilege
  - Can be authenticated internally or externally(using OS)
- User can be grouped
- Roles are collection of privilege
  - Some dbms have predefined or you can have user-defined roles.
- Create roles or groups, assign new member to group, not each
- Principle of least privilege
- Create granular groups

## Managing access

- Privilege can be granted to user, groups or roles.

```
GRANT CONNECT TO 'salesteam'; -- Sales team is group
-- OR you can allow a specific user
REVOKE CONNECT TO 'public';
GRANT SELECT ON mydb.table TO 'john';
GRANT CREATE TABLE TO 'salesteam';

GRANT VIEW ON mydb.myproc to 'salesteam';
-- You can also grant EXECUTE, ALTER to procedures

REVOKE SELECT ON mydb.mytable TO 'salesteam';
-- After revoke another role can be used by team member to do that
DENY SELECT ON mydb.mytable TO 'salesteam';
-- After deny they can't do that. i.e. it overrides an existing granted
permission.
```

## Auditing db activities

Doesn't directly prevent unauthorized access but helps you identify gaps in security measures and track errors.

- Auditing = recording activity (logs) and reviewing them
- Track and review activity (failed or successful)
- Trigger can be used for logging

## Encrypting data

Adds a layer of security, last line of defence

- May be required.

- Data at rest should be encrypted (Transparent data encryption) to ensure that malicious users cant use the files.
  - DB level, table level, column (but performance and complexity)
- Symmetric: same key for encryption or decryption
  - Key gone , data gone , your job done
- Assymetric: one public key, one private key.
- Transparent data encryption is not visible to user. Automated.
- Customer managed key => Bring your own key(BYOK)
  - Even db admins cannot access data
- Encryption takes time.
- Assymetric have greater overheads
  - So, mostly symmetric is enough
- Often provided by RDBMS
  - TLS : Transport layer security
  - SSL : Secure socket level
  - Allows enable/disable